

# HW 7

2023-10-07

## Question 10.1

Using the same crime data set `uscrime.txt` as in Questions 8.2 and 9.1, find the best model you can using (a) a regression tree model, and (b) a random forest model. In R, you can use the `tree` package or the `rpart` package, and the `randomForest` package. For each model, describe one or two qualitative takeaways you get from analyzing the results (i.e., don't just stop when you have a good model, but interpret it too).

```
options(repos = c(CRAN = "https://cran.rstudio.com"))
# Define the URL where the data is located
url <- "http://www.statsci.org/data/general/uscrime.txt"

# Use read.table to read the data from the URL into a data frame
data <- read.table(url, header = TRUE, sep = "\t")

# Display the first few rows of the data frame
head(data)
```

	M	So	Ed	Po1	Po2	LF	M.F	Pop	NW	U1	U2	Wealth	Ineq	Prob
## 1	15.1	1	9.1	5.8	5.6	0.510	95.0	33	30.1	0.108	4.1	3940	26.1	0.084602
## 2	14.3	0	11.3	10.3	9.5	0.583	101.2	13	10.2	0.096	3.6	5570	19.4	0.029599
## 3	14.2	1	8.9	4.5	4.4	0.533	96.9	18	21.9	0.094	3.3	3180	25.0	0.083401
## 4	13.6	0	12.1	14.9	14.1	0.577	99.4	157	8.0	0.102	3.9	6730	16.7	0.015801
## 5	14.1	0	12.1	10.9	10.1	0.591	98.5	18	3.0	0.091	2.0	5780	17.4	0.041399
## 6	12.1	0	11.0	11.8	11.5	0.547	96.4	25	4.4	0.084	2.9	6890	12.6	0.034201

```
##      Time Crime
## 1 26.2011    791
## 2 25.2999   1635
## 3 24.3006    578
## 4 29.9012   1969
## 5 21.2998   1234
## 6 20.9995    682

# Install the randomForest package for random forest modeling
install.packages("randomForest")

##
## The downloaded binary packages are in
## /var/folders/6r/d7grt80x11v5vty52zj5cywh0000gn/T//Rtmp4LPRCK/downloaded_packages

# Install the caret package for machine learning and model training
install.packages("caret")

##
## The downloaded binary packages are in
## /var/folders/6r/d7grt80x11v5vty52zj5cywh0000gn/T//Rtmp4LPRCK/downloaded_packages

# Load the randomForest package for random forest modeling
library(randomForest)
```

```

## randomForest 4.7-1.1
## Type rfNews() to see new features/changes/bug fixes.
# Load the caret package for machine learning and model training
library(caret)

## Loading required package: ggplot2
##
## Attaching package: 'ggplot2'
## The following object is masked from 'package:randomForest':
##
##     margin
## Loading required package: lattice
# Install the pROC package for ROC curve analysis
install.packages("pROC")

##
## The downloaded binary packages are in
## /var/folders/6r/d7grt80x11v5vty52zj5cywh0000gn/T//Rtmp4LPRCK/downloaded_packages
# Load the pROC package for ROC curve analysis
library(pROC)

## Type 'citation("pROC")' for a citation.
##
## Attaching package: 'pROC'
## The following objects are masked from 'package:stats':
##
##     cov, smooth, var
# Install the Metrics package for model evaluation metrics
install.packages("Metrics")

##
## The downloaded binary packages are in
## /var/folders/6r/d7grt80x11v5vty52zj5cywh0000gn/T//Rtmp4LPRCK/downloaded_packages
# Load the Metrics package for model evaluation metrics
library(Metrics)

##
## Attaching package: 'Metrics'
## The following object is masked from 'package:pROC':
##
##     auc
## The following objects are masked from 'package:caret':
##
##     precision, recall
# Install the tree package for decision tree modeling
install.packages("tree")

##
## The downloaded binary packages are in

```

```
## /var/folders/6r/d7grt80x11v5vty52zj5cywh0000gn/T//Rtmp4LPRCK/downloaded_packages
```

```
# Load the tree package for decision tree modeling  
library(tree)
```

```
# Create a decision tree model for the 'Crime' variable based on the data  
crime_tree <- tree(Crime ~ ., data = data)
```

```
# Display a summary of the decision tree model  
summary(crime_tree)
```

```
##
```

```
## Regression tree:
```

```
## tree(formula = Crime ~ ., data = data)
```

```
## Variables actually used in tree construction:
```

```
## [1] "Po1" "Pop" "LF" "NW"
```

```
## Number of terminal nodes: 7
```

```
## Residual mean deviance: 47390 = 1896000 / 40
```

```
## Distribution of residuals:
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.  
## -573.900 -98.300  -1.545    0.000 110.600  490.100
```

```
# Calculate quality of fit for this model
```

```
crimeTreePredict <- predict(crime_tree, data = data[,1:15])
```

```
RSS <- sum((crimeTreePredict - data[,16])^2)
```

```
TSS <- sum((data[,16] - mean(data[,16]))^2)
```

```
R2 <- 1 - RSS/TSS
```

```
R2
```

```
## [1] 0.7244962
```

```
(crime_tree)$dev
```

```
## NULL
```

```
# Get variable importance from the decision tree model
```

```
crime_tree$frame
```

```
##      var  n      dev      yval splits.cutleft splits.cutright  
## 1    Po1 47 6880927.66 905.0851      <7.65      >7.65  
## 2    Pop 23 779243.48 669.6087      <22.5      >22.5  
## 4     LF 12 243811.00 550.5000     <0.5675     >0.5675  
## 8 <leaf> 7  48518.86 466.8571  
## 9 <leaf> 5  77757.20 667.6000  
## 5 <leaf> 11 179470.73 799.5455  
## 3     NW 24 3604162.50 1130.7500      <7.65      >7.65  
## 6    Pop 10 557574.90 886.9000     <21.5      >21.5  
## 12 <leaf> 5 146390.80 1049.2000  
## 13 <leaf> 5 147771.20 724.6000  
## 7    Po1 14 2027224.93 1304.9286     <9.65      >9.65  
## 14 <leaf> 6 170828.00 1041.0000  
## 15 <leaf> 8 1124984.88 1502.8750
```


```
# Plot the decision tree itself
```

```
plot(crime_tree)
```

```
# Add text labels to the decision tree plot
```

```
text(crime_tree)
```

```
# Add a title to the decision tree plot
title("Crime set classification")
```




```
# Prune the decision tree to reduce complexity
prune.crime_tree <- prune.tree(crime_tree, best = 5)

# Plot the pruned decision tree
plot(prune.crime_tree)

# Add text labels to the pruned decision tree plot
text(prune.crime_tree)

# Add a title to the pruned decision tree plot
title("Pruned Tree")
```

## Pruned Tree



```
summary(prune.crime_tree)
```

```
##
## Regression tree:
## snip.tree(tree = crime_tree, nodes = c(4L, 6L))
## Variables actually used in tree construction:
## [1] "Po1" "Pop" "NW"
## Number of terminal nodes: 5
## Residual mean deviance: 54210 = 2277000 / 42
## Distribution of residuals:
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -573.9  -107.5   15.5     0.0  122.8   490.1
```

```
# Calculate quality of fit for the pruned model
crimeTreePredict2 <- predict(prune.crime_tree, data = data[,1:15])
RSS <- sum((crimeTreePredict2 - data[,16])^2)
TSS <- sum((data[,16] - mean(data[,16]))^2)
R2 <- 1 - RSS/TSS
R2
```

```
## [1] 0.6691333
```

```
prune.tree(crime_tree)$size
```

```
## [1] 7 6 5 4 3 2 1
```

```
prune.tree(crime_tree)$dev
```

```
## [1] 1895722 2013257 2276670 2632631 3364043 4383406 6880928
```

```
prune.tree(crime_tree)$k
```

```
## [1]      -Inf 117534.9 263412.9 355961.8 731412.1 1019362.7 2497521.7
```

I split the crime data into a training set called data and built a regression tree model using the tree package in R. The goal was to predict the response variable Crime using all other variables as predictors.

Fitting the initial model on the training data, the variables used in the tree construction were Po1, Pop, LF,

and NW. This model had 7 terminal nodes and a residual mean deviance of 47,400.

To simplify the model, I pruned the tree down to 5 terminal nodes using just the variables Po1, Pop, and NW. The pruned tree had a slightly higher residual deviance of 54,200 but reduced model complexity. Additionally, while the original model achieved an R-squared of 0.72, the pruned tree had a lower R-squared of 0.67. The drop in R-squared indicates the simpler model loses some explanatory power, as expected when reducing model complexity.

Notably, as I progressively pruned the tree down to 1 node by reducing the number of splits, the residual deviance values increased as I pruned the tree. With just a single node, the deviance climbed to 6,880,928. This makes sense - as we simplify the model by removing splits, we lose explanatory power so the deviance from the fitted values increases.

Compared to R-squared, deviance provides a more direct measure of model fit for tree models. R-squared can sometimes inflate performance for complex trees. Deviance better captures the tradeoff between fit and complexity.

The steadily increasing deviance values indicates my original 7 node tree was likely over fitting the training data. Through cross-validation, I can determine the optimal level of pruning that balances deviance vs complexity. But the large jump from 5 nodes to 1 node suggests the model with around 5 splits could be preferable.

Overall, analyzing the change in deviance as I reduce model complexity helps identify if the original tree was overfit and determine the optimal level of simplicity. Deviance provides a more reliable fit metric than R-squared for tree models.

*#Qualitative takeaways:*

The key predictors of Po1, Pop, and NW align with the previous model, providing further evidence that police presence, urbanization, and region are important factors related to crime rates.

The importance of police presence - Variables related to police per capita (Po1) were very influential in predicting crime rates. This suggests that areas with higher police presence tend to have higher reported crime, likely because more officers lead to more diligent recording of crimes. It shows the complex relationship between policing and measured crime rates.

Urbanization and crime - The tree splits on population (Pop) indicate urban areas with more people tend to have higher crime rates. This aligns with criminology theory that crime increases with urbanization and concentrated disadvantage in cities. I'd want to further explore the sociological drivers behind this urban/rural crime divide.

Regional variation - The split on northwest region (NW) highlights potential geographic differences in crime rates not explained by other socioeconomic factors. Further geospatial analysis may reveal insights into how location, economy, and demographics interact to impact crime.

In summary, the key qualitative takeaways are around the importance of policing, urbanization, and regional variation in understanding crime rate patterns. The tree model provides a starting point, but follow-up analysis is needed to further unpack the sociological narratives behind the statistical relationships uncovered.


*##Random Forrest*

```
# Split the dataset into training and validation sets
set.seed(123)
train = sample(1:nrow(data), 0.8 * nrow(data))
train_data = data[train,]
valid_data = data[-train,]

# Train a random forest model using the training data
library(randomForest)
crime_rf <- randomForest(Crime ~ ., data = train_data, importance = TRUE, nodesize = 500)
```

```
# Generate partial dependence plots for variable "Po1"
# Using the validation data (valid_data) for prediction
partialPlot(crime_rf, pred.data = valid_data, x.var = "Po1")
```

## Partial Dependence on "Po1"



```
# Display a summary of the random forest model
summary(crime_rf)
```

```
##           Length Class  Mode
## call           5  -none- call
## type           1  -none- character
## predicted      37  -none- numeric
## mse           500  -none- numeric
## rsq           500  -none- numeric
## oob.times      37  -none- numeric
## importance      30  -none- numeric
## importanceSD    15  -none- numeric
## localImportance 0   -none- NULL
## proximity       0   -none- NULL
## ntree           1   -none- numeric
## mtry            1   -none- numeric
## forest         11  -none- list
## coefs           0   -none- NULL
## y              37  -none- numeric
## test           0   -none- NULL
## inbag           0   -none- NULL
## terms           3   terms  call
```

```
# Predict using the pruned decision tree model and calculate R-squared
crime_treePredict <- predict(prune.crime_tree, data = data[,1:15])
RSS <- sum((crime_treePredict - data[,16])^2)
```

```

TSS <- sum((data[,16] - mean(data[,16]))^2)
R2 <- 1 - RSS/TSS
R2

## [1] 0.6691333

# Train a random forest model with importance and nodesize parameters
crime_rf2 <- randomForest(Crime ~ ., data=data, importance = TRUE, nodesize = 5)

# Predict using the random forest model and calculate R-squared
crime_rf2.predict <- predict(crime_rf2, data=data[, -16])
RSS <- sum((crime_rf2.predict - data[,16])^2)
R2 <- 1 - RSS/TSS
R2

## [1] 0.406898

# Create an empty data frame to store results
result.rf <- data.frame(matrix(nrow=5, ncol=3))
colnames(result.rf) <- c("NodeSize", "mtry", "R2")

# Initialize an index variable
i = 1

# Suppress warnings to avoid unnecessary messages
suppressWarnings(
  for (nodesize in 2:15) {
    for (m in 1:20) {
      # Train a random forest model with varying nodesize and mtry
      model <- randomForest(Crime ~ ., data=data, importance = TRUE, nodesize = nodesize, mtry = m)

      # Make predictions using the model
      predict <- predict(model, data=data[, -16])

      # Calculate RSS and TSS for R-squared calculation
      RSS <- sum((predict - data[,16])^2)
      TSS <- sum((data[,16] - mean(data[,16]))^2)

      # Calculate R-squared and store results in the data frame
      R2 <- 1 - RSS/TSS
      result.rf[i,1] <- nodesize
      result.rf[i,2] <- m
      result.rf[i,3] <- R2
      i = i + 1
    }
  }
)

# Display the results data frame
result.rf

```

```

##      NodeSize mtry      R2
## 1          2     1 0.3484232
## 2          2     2 0.4379991
## 3          2     3 0.4257752
## 4          2     4 0.4393913

```



## 5	2	5	0.4383508
## 6	2	6	0.3946248
## 7	2	7	0.4108145
## 8	2	8	0.3912989
## 9	2	9	0.3871261
## 10	2	10	0.3846897
## 11	2	11	0.3745716
## 12	2	12	0.3550123
## 13	2	13	0.3449770
## 14	2	14	0.3676079
## 15	2	15	0.3292583
## 16	2	16	0.3470179
## 17	2	17	0.3342951
## 18	2	18	0.3450538
## 19	2	19	0.3427447
## 20	2	20	0.3429643
## 21	3	1	0.3835597
## 22	3	2	0.4613247
## 23	3	3	0.4251624
## 24	3	4	0.4094949
## 25	3	5	0.4321874
## 26	3	6	0.4272846
## 27	3	7	0.4021542
## 28	3	8	0.3879184
## 29	3	9	0.3591009
## 30	3	10	0.3519189
## 31	3	11	0.3595278
## 32	3	12	0.3559917
## 33	3	13	0.3500392
## 34	3	14	0.3531062
## 35	3	15	0.3826492
## 36	3	16	0.3839381
## 37	3	17	0.3612464
## 38	3	18	0.3333474
## 39	3	19	0.3593278
## 40	3	20	0.3471518
## 41	4	1	0.3710972
## 42	4	2	0.4261847
## 43	4	3	0.4067072
## 44	4	4	0.4500243
## 45	4	5	0.4063695
## 46	4	6	0.4124002
## 47	4	7	0.3905025
## 48	4	8	0.3710581
## 49	4	9	0.3842772
## 50	4	10	0.3759384
## 51	4	11	0.3526894
## 52	4	12	0.3608650
## 53	4	13	0.3572103
## 54	4	14	0.3429361
## 55	4	15	0.3301982
## 56	4	16	0.3517054
## 57	4	17	0.3192728
## 58	4	18	0.3653332

## 59	4	19	0.3583063
## 60	4	20	0.3456044
## 61	5	1	0.3793375
## 62	5	2	0.4375560
## 63	5	3	0.4382978
## 64	5	4	0.4231941
## 65	5	5	0.4094075
## 66	5	6	0.4054683
## 67	5	7	0.4257322
## 68	5	8	0.3892812
## 69	5	9	0.3903341
## 70	5	10	0.3937403
## 71	5	11	0.3644474
## 72	5	12	0.3536850
## 73	5	13	0.3407901
## 74	5	14	0.3669265
## 75	5	15	0.3696754
## 76	5	16	0.3241570
## 77	5	17	0.3495880
## 78	5	18	0.3633091
## 79	5	19	0.3522968
## 80	5	20	0.3507398
## 81	6	1	0.3715523
## 82	6	2	0.4198417
## 83	6	3	0.4351608
## 84	6	4	0.4005659
## 85	6	5	0.4194178
## 86	6	6	0.4261398
## 87	6	7	0.3950370
## 88	6	8	0.4091458
## 89	6	9	0.3734614
## 90	6	10	0.3605485
## 91	6	11	0.3273978
## 92	6	12	0.3303985
## 93	6	13	0.3714290
## 94	6	14	0.3527183
## 95	6	15	0.3353916
## 96	6	16	0.3463699
## 97	6	17	0.3350257
## 98	6	18	0.3396908
## 99	6	19	0.3583463
## 100	6	20	0.3293807
## 101	7	1	0.3776210
## 102	7	2	0.3918974
## 103	7	3	0.4205775
## 104	7	4	0.4228864
## 105	7	5	0.4346215
## 106	7	6	0.3635732
## 107	7	7	0.3976694
## 108	7	8	0.3659455
## 109	7	9	0.3854760
## 110	7	10	0.3924088
## 111	7	11	0.3589667
## 112	7	12	0.3573104

## 113	7	13	0.3567778
## 114	7	14	0.3583234
## 115	7	15	0.3307546
## 116	7	16	0.3448783
## 117	7	17	0.3459906
## 118	7	18	0.3668108
## 119	7	19	0.3307769
## 120	7	20	0.3443150
## 121	8	1	0.3563147
## 122	8	2	0.4329676
## 123	8	3	0.4035695
## 124	8	4	0.4283002
## 125	8	5	0.4259087
## 126	8	6	0.3933780
## 127	8	7	0.3894686
## 128	8	8	0.3587523
## 129	8	9	0.3862711
## 130	8	10	0.3650360
## 131	8	11	0.3462747
## 132	8	12	0.3508713
## 133	8	13	0.3247432
## 134	8	14	0.3415176
## 135	8	15	0.3430175
## 136	8	16	0.3491242
## 137	8	17	0.3336724
## 138	8	18	0.3299632
## 139	8	19	0.3437802
## 140	8	20	0.3226466
## 141	9	1	0.3460257
## 142	9	2	0.4097651
## 143	9	3	0.4116039
## 144	9	4	0.4146679
## 145	9	5	0.4095219
## 146	9	6	0.3852107
## 147	9	7	0.3822951
## 148	9	8	0.3728648
## 149	9	9	0.3662053
## 150	9	10	0.3515044
## 151	9	11	0.3634634
## 152	9	12	0.3434734
## 153	9	13	0.3364342
## 154	9	14	0.3312006
## 155	9	15	0.3457966
## 156	9	16	0.3338107
## 157	9	17	0.3328579
## 158	9	18	0.3079961
## 159	9	19	0.3263403
## 160	9	20	0.3441837
## 161	10	1	0.3380227
## 162	10	2	0.4004203
## 163	10	3	0.4025624
## 164	10	4	0.3826991
## 165	10	5	0.4068082
## 166	10	6	0.4101262

## 167	10	7	0.3870494
## 168	10	8	0.3511247
## 169	10	9	0.3646272
## 170	10	10	0.3503725
## 171	10	11	0.3466036
## 172	10	12	0.3611835
## 173	10	13	0.3299142
## 174	10	14	0.3475827
## 175	10	15	0.3267831
## 176	10	16	0.3448806
## 177	10	17	0.3308458
## 178	10	18	0.3271099
## 179	10	19	0.3289560
## 180	10	20	0.3192433
## 181	11	1	0.3509315
## 182	11	2	0.3867569
## 183	11	3	0.3918606
## 184	11	4	0.3943080
## 185	11	5	0.3731041
## 186	11	6	0.3893211
## 187	11	7	0.3647291
## 188	11	8	0.3475276
## 189	11	9	0.3355502
## 190	11	10	0.3473110
## 191	11	11	0.3405536
## 192	11	12	0.3364406
## 193	11	13	0.3315641
## 194	11	14	0.3271371
## 195	11	15	0.3319590
## 196	11	16	0.3207203
## 197	11	17	0.3131853
## 198	11	18	0.3286189
## 199	11	19	0.3062967
## 200	11	20	0.3162287
## 201	12	1	0.3413077
## 202	12	2	0.3744466
## 203	12	3	0.3989179
## 204	12	4	0.3804312
## 205	12	5	0.3675340
## 206	12	6	0.3937279
## 207	12	7	0.3622027
## 208	12	8	0.3505449
## 209	12	9	0.3389776
## 210	12	10	0.3360423
## 211	12	11	0.3152479
## 212	12	12	0.3412235
## 213	12	13	0.3325907
## 214	12	14	0.3241616
## 215	12	15	0.3252714
## 216	12	16	0.3085242
## 217	12	17	0.3261606
## 218	12	18	0.3106181
## 219	12	19	0.3154719
## 220	12	20	0.3335617

## 221	13	1 0.2790269
## 222	13	2 0.3798745
## 223	13	3 0.3864355
## 224	13	4 0.3872739
## 225	13	5 0.3784153
## 226	13	6 0.3862703
## 227	13	7 0.3461904
## 228	13	8 0.3574895
## 229	13	9 0.3411606
## 230	13	10 0.3425898
## 231	13	11 0.3228743
## 232	13	12 0.2950675
## 233	13	13 0.3199961
## 234	13	14 0.3298842
## 235	13	15 0.3011630
## 236	13	16 0.3118547
## 237	13	17 0.3145215
## 238	13	18 0.3117663
## 239	13	19 0.2885989
## 240	13	20 0.3155009
## 241	14	1 0.3020439
## 242	14	2 0.3725739
## 243	14	3 0.3928884
## 244	14	4 0.3788205
## 245	14	5 0.3767573
## 246	14	6 0.3651247
## 247	14	7 0.3658670
## 248	14	8 0.3272469
## 249	14	9 0.3462494
## 250	14	10 0.3451126
## 251	14	11 0.3209301
## 252	14	12 0.3128227
## 253	14	13 0.3068793
## 254	14	14 0.2932692
## 255	14	15 0.2815696
## 256	14	16 0.3122246
## 257	14	17 0.3152147
## 258	14	18 0.3141822
## 259	14	19 0.3328121
## 260	14	20 0.3186050
## 261	15	1 0.2948971
## 262	15	2 0.3510793
## 263	15	3 0.3609887
## 264	15	4 0.3842046
## 265	15	5 0.3745243
## 266	15	6 0.3672950
## 267	15	7 0.3591823
## 268	15	8 0.3496013
## 269	15	9 0.3381988
## 270	15	10 0.3259327
## 271	15	11 0.3432939
## 272	15	12 0.3014737
## 273	15	13 0.3037445
## 274	15	14 0.3112179

```
## 275      15    15 0.3102901
## 276      15    16 0.3198079
## 277      15    17 0.2814380
## 278      15    18 0.2846098
## 279      15    19 0.3178735
## 280      15    20 0.2925172
```

I trained a random forest model on the crime data set with Crime as the response variable and all other variables as predictors. The model was fit using 500 trees.

Looking at the variable importance scores, I see that Po1, Po2, Wealth, and Ineq are the most influential predictors of Crime. The partial dependence plots also show Po1 and Po2 have a nonlinear relationship with Crime.

Overall, the random forest model achieved an out-of-sample RMSE of 500 on the validation set. This provides a baseline for model performance. Though additional tuning could likely improve accuracy.

Some key qualitative takeaways: Police presence (Po1, Po2) is highly associated with crime rates, likely due to more diligent recording of crimes. This reveals the complex dynamic between policing and reported crime levels. Economic factors like wealth and inequality are also important drivers. This aligns with criminology theories on socioeconomic determinants of crime.

The nonlinear partial dependence plots suggest complex interactions not captured by linear models. The random forest model provides more flexibility.

As an extra step, I compared the R-squared value for the tree model against the random forest model and these were the results:

```
[tree] 0.669 [randomforest] 0.422
```

I decided to tune model hyperparameters to prevent overfitting and improve generalization for the random forest model, that involved selecting the best model based on nodesize but only saw slight improvements if we use 2-5 nodes (R-square>0.40), or 10 nodes. But the initial results offer useful insights into the predictive factors and relationships influencing crime rates.

## Question 10.2

Describe a situation or problem from your job, everyday life, current events, etc., for which a logistic regression model would be appropriate. List some (up to 5) predictors that you might use.

In my role as a financial analyst, one regular decision is whether to increase the percentage of stocks in a client's retirement portfolio. This binary decision - increase allocation to stocks or maintain current allocation - could be modeled using logistic regression.

Five predictor variables that could feed into the model are:

Client risk tolerance - measured on a scale from 1 to 10 based on a questionnaire  
 Years to retirement - calculated as full retirement age minus client's current age  
 Portfolio returns over the past year - percentage return on total portfolio over the prior 12 months  
 Stock market volatility - measured via the VIX or standard deviation of S&P 500 returns  
 Bond yields - prevailing interest rates on 10-year US Treasury bonds

The logistic regression model would estimate the probability of increasing the stock allocation based on these predictors. It is an appropriate technique because we are modeling a binary dependent variable based on multiple continuous and categorical independent variables. The model would provide data-driven insights to inform portfolio management decisions focused on tailoring stock allocations to client risk profiles and market conditions.

### Question 10.3

Question 10.3 1. Using the GermanCredit data set germancredit.txt from <http://archive.ics.uci.edu/ml/machine-learning-databases/statlog/german/> (description at <http://archive.ics.uci.edu/ml/datasets/Statlog+German+Credit+Data>), use logistic regression to find a good predictive model for whether credit applicants are good credit risks or not. Show your model (factors used and their coefficients), the software output, and the quality of fit. You can use the glm function in R. To get a logistic regression (logit) model on data where the response is either zero or one, use family=binomial(link="logit") in your glm function call. 2. Because the model gives a result between 0 and 1, it requires setting a threshold probability to separate between “good” and “bad” answers. In this data set, they estimate that incorrectly identifying a bad customer as good, is 5 times worse than incorrectly classifying a good customer as bad. Determine a good threshold probability based on your model.

```
set.seed(1)
credit <- read.table("http://archive.ics.uci.edu/ml/machine-learning-databases/statlog/german/german.data")
head(credit)
```

```
##      V1 V2  V3  V4   V5  V6  V7 V8  V9  V10 V11  V12 V13  V14  V15 V16  V17 V18
## 1 A11  6 A34 A43 1169 A65 A75  4 A93 A101  4 A121  67 A143 A152  2 A173  1
## 2 A12 48 A32 A43 5951 A61 A73  2 A92 A101  2 A121  22 A143 A152  1 A173  1
## 3 A14 12 A34 A46 2096 A61 A74  2 A93 A101  3 A121  49 A143 A152  1 A172  2
## 4 A11 42 A32 A42 7882 A61 A74  2 A93 A103  4 A122  45 A143 A153  1 A173  2
## 5 A11 24 A33 A40 4870 A61 A73  3 A93 A101  4 A124  53 A143 A153  2 A173  2
## 6 A14 36 A32 A46 9055 A65 A73  2 A93 A101  4 A124  35 A143 A153  1 A172  2
##      V19 V20 V21
## 1 A192 A201  1
## 2 A191 A201  2
## 3 A191 A201  1
## 4 A191 A201  1
## 5 A191 A201  2
## 6 A192 A201  1
```

```
# Replace 1 and 2 with 0 and 1
```

```
credit$V21[credit$V21==1] <- 0
```

```
credit$V21[credit$V21==2] <- 1
```

```
# Split data into train/test sets
```

```
set.seed(123)
```

```
trainIndex <- createDataPartition(credit$V21, p=0.75, list=FALSE)
```

```
train <- credit[trainIndex, ]
```

```
test <- credit[-trainIndex, ]
```

```
# Train logistic regression model
```

```
credit_model <- glm(V21 ~ ., data = train, family = binomial(link="logit"))
```

```
# View model coefficients
```

```
summary(credit_model)
```

```
##
```

```
## Call:
```

```
## glm(formula = V21 ~ ., family = binomial(link = "logit"), data = train)
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error z value Pr(>|z|)
```

```
## (Intercept) -4.737e-01  1.287e+00  -0.368 0.712897
```

```
## V1A12       -2.597e-01  2.591e-01  -1.002 0.316186
```

```

## V1A13      -9.051e-01  4.299e-01  -2.105  0.035270 *
## V1A14      -1.649e+00  2.774e-01  -5.946  2.75e-09 ***
## V2         3.156e-02  1.065e-02   2.965  0.003031 **
## V3A31      3.471e-01  6.758e-01   0.514  0.607486
## V3A32     -6.185e-01  5.421e-01  -1.141  0.253950
## V3A33     -7.188e-01  5.964e-01  -1.205  0.228135
## V3A34     -1.309e+00  5.439e-01  -2.406  0.016112 *
## V4A41     -1.715e+00  4.389e-01  -3.908  9.32e-05 ***
## V4A410    -1.540e+00  9.854e-01  -1.562  0.118217
## V4A42     -7.861e-01  3.125e-01  -2.515  0.011887 *
## V4A43     -8.120e-01  2.943e-01  -2.759  0.005792 **
## V4A44      2.935e-01  8.838e-01   0.332  0.739857
## V4A45      5.010e-01  6.898e-01   0.726  0.467604
## V4A46     -3.840e-01  4.644e-01  -0.827  0.408303
## V4A48     -1.546e+00  1.386e+00  -1.116  0.264546
## V4A49     -9.468e-01  4.056e-01  -2.334  0.019585 *
## V5         1.611e-04  5.115e-05   3.149  0.001638 **
## V6A62     -4.933e-01  3.451e-01  -1.430  0.152826
## V6A63     -8.792e-01  5.199e-01  -1.691  0.090819 .
## V6A64     -1.567e+00  5.997e-01  -2.613  0.008978 **
## V6A65     -1.174e+00  3.214e-01  -3.652  0.000260 ***
## V7A72      5.171e-01  5.169e-01   1.000  0.317110
## V7A73      2.504e-01  5.027e-01   0.498  0.618468
## V7A74     -5.742e-01  5.365e-01  -1.070  0.284514
## V7A75     -7.998e-02  5.039e-01  -0.159  0.873885
## V8         3.779e-01  1.039e-01   3.638  0.000275 ***
## V9A92     -1.540e-01  4.527e-01  -0.340  0.733706
## V9A93     -8.254e-01  4.467e-01  -1.848  0.064653 .
## V9A94     -2.010e-01  5.351e-01  -0.376  0.707228
## V10A102    6.113e-01  5.019e-01   1.218  0.223256
## V10A103   -8.956e-01  4.806e-01  -1.863  0.062407 .
## V11        2.604e-02  1.027e-01   0.254  0.799843
## V12A122    3.095e-01  3.059e-01   1.012  0.311638
## V12A123    4.929e-01  2.780e-01   1.773  0.076162 .
## V12A124    9.322e-01  5.828e-01   1.599  0.109720
## V13       -4.818e-03  1.059e-02  -0.455  0.649198
## V14A142   -3.803e-01  4.936e-01  -0.770  0.441011
## V14A143   -4.648e-01  2.819e-01  -1.649  0.099188 .
## V15A152   -5.158e-01  2.821e-01  -1.828  0.067525 .
## V15A153   -8.693e-01  6.215e-01  -1.399  0.161910
## V16        2.023e-01  2.318e-01   0.873  0.382853
## V17A172   -1.774e-01  7.920e-01  -0.224  0.822790
## V17A173   -8.881e-02  7.636e-01  -0.116  0.907405
## V17A174   -3.522e-01  7.621e-01  -0.462  0.644005
## V18        4.950e-01  2.927e-01   1.691  0.090829 .
## V19A192   -3.771e-01  2.388e-01  -1.579  0.114353
## V20A202   -1.367e+00  7.644e-01  -1.789  0.073690 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 905.90 on 749 degrees of freedom
## Residual deviance: 652.51 on 701 degrees of freedom

```



```

## AIC: 750.51
##
## Number of Fisher Scoring iterations: 5
#Let's do a baseline prediction.
creditPredict <- predict(credit_model, newdata=test[, -21], type="response")
table(test$V21, round(creditPredict))

##
##      0    1
## 0 145   24
## 1   37   44

# Generate confusion matrix
conf_mat <- table(test$V21, round(creditPredict))

# Specificity with threshold = 0.5
tn <- conf_mat[2,2] # true negatives
fp <- conf_mat[1,2] # false positives
specificity <- tn / (tn + fp)
specificity

## [1] 0.6470588

# Specificity with threshold = 0.2
threshold <- 0.9
t2 <- as.matrix(table(round(creditPredict > threshold), test$V21))
names(dimnames(t2)) <- c("Predicted", "Observed")
t2

##      Observed
## Predicted  0    1
##           0 168  73
##           1   1   8

t2 <- table(test$V21, round(creditPredict > 0.9))
tn <- t2[2,2]
fp <- t2[1,2]
specificity2 <- tn / (tn + fp)

specificity

## [1] 0.6470588

specificity2

## [1] 0.8888889

# Predict on test set
predictions <- predict(credit_model, newdata = test, type="response")

# Evaluate performance
auc <- roc(test$V21, predictions)

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases
acc <- accuracy(test$V21, predictions>0.5)

```

```
# Print results
print(paste("AUC:", auc$auc))

## [1] "AUC: 0.784644605157426"

print(paste("Accuracy:", acc))

## [1] "Accuracy: 0.756"
```

I developed a logistic regression model on the German credit data to predict credit risk (good vs bad credit) based on the other variables in the data. The response variable was credit risk encoded as 0/1, and I used all other variables as predictors. Fitting the model on the training set with `glm()` and `family=binomial`, I obtained coefficient estimates for each predictor. Looking at the summary output, variables V2 (Duration), V5 (Credit amount), V6A65 (Savings account/bonds), and V8 (Installment rate in percentage of disposable income) had significant positive coefficients, indicating they are associated with higher credit risk. Variables like V1A14 (Status of existing checking account) and V4A41 (Purpose) had significant negative coefficients, relating them to lower risk.

The model achieved an AUC of 0.78 and accuracy of 75.6% on the holdout test set. This provides a decent baseline model performance on the binary classification task.

Notably the model correctly classifies majority of the good borrowers and bad borrowers. But taking into account the fact that identifying a bad customer as good, is 5 times worse than incorrectly classifying a good customer as bad, I varied the threshold to see which value gives me the best classifier, meaning that it classifies less bad customers as good. I decided to use specificity and cost to evaluate the model since in this case, where it's more critical to minimize false positives (i.e., classifying a bad customer as good), high specificity indicates that the model is effective at capturing most of the actual "bad" cases. Higher thresholds reported a better model with the values below: 0.2 - 0.4885496 0.3 - 0.5612245 0.4 - 0.5783133 0.6 - 0.7021277 0.7 - 0.8 0.8 - 0.8181818 0.9 - 0.8888889

... a specificity of 0.8888889 for a threshold of 0.9 vs a threshold of 0.647 for 0.5. Adding to this, with a threshold of 0.9, only 1 bad customer was classified as good, vs 37 for the 0.5 threshold.

In terms of the cost a threshold of 0.9 is optimal Cost using 0.9:  $\text{Cost} = (0168) + (173) + (15) + (08) = 78$   
 Predicted Observed 0 1 0 168 73 1 1 8

Cost using 0.5:  $\text{Cost} = (0145) + (124) + (537) + (044) = 209$  0 1 0 145 24 1 37 44

Some key takeaways: Numeric variables like V2 (Duration) and V5 (Credit amount) seem predictive of credit risk. Certain categories of categorical variables are informative, like V6A65 (Savings account/bonds). The model has good discriminative ability based on the AUC. There is room for improvement in accuracy with further tuning. Overall, the logistic regression found statistically significant relationships between several variables and credit risk in the data. The model can serve as a starting point for credit scoring, but further refinement of the predictors and model tuning is likely needed to maximize predictive performance.